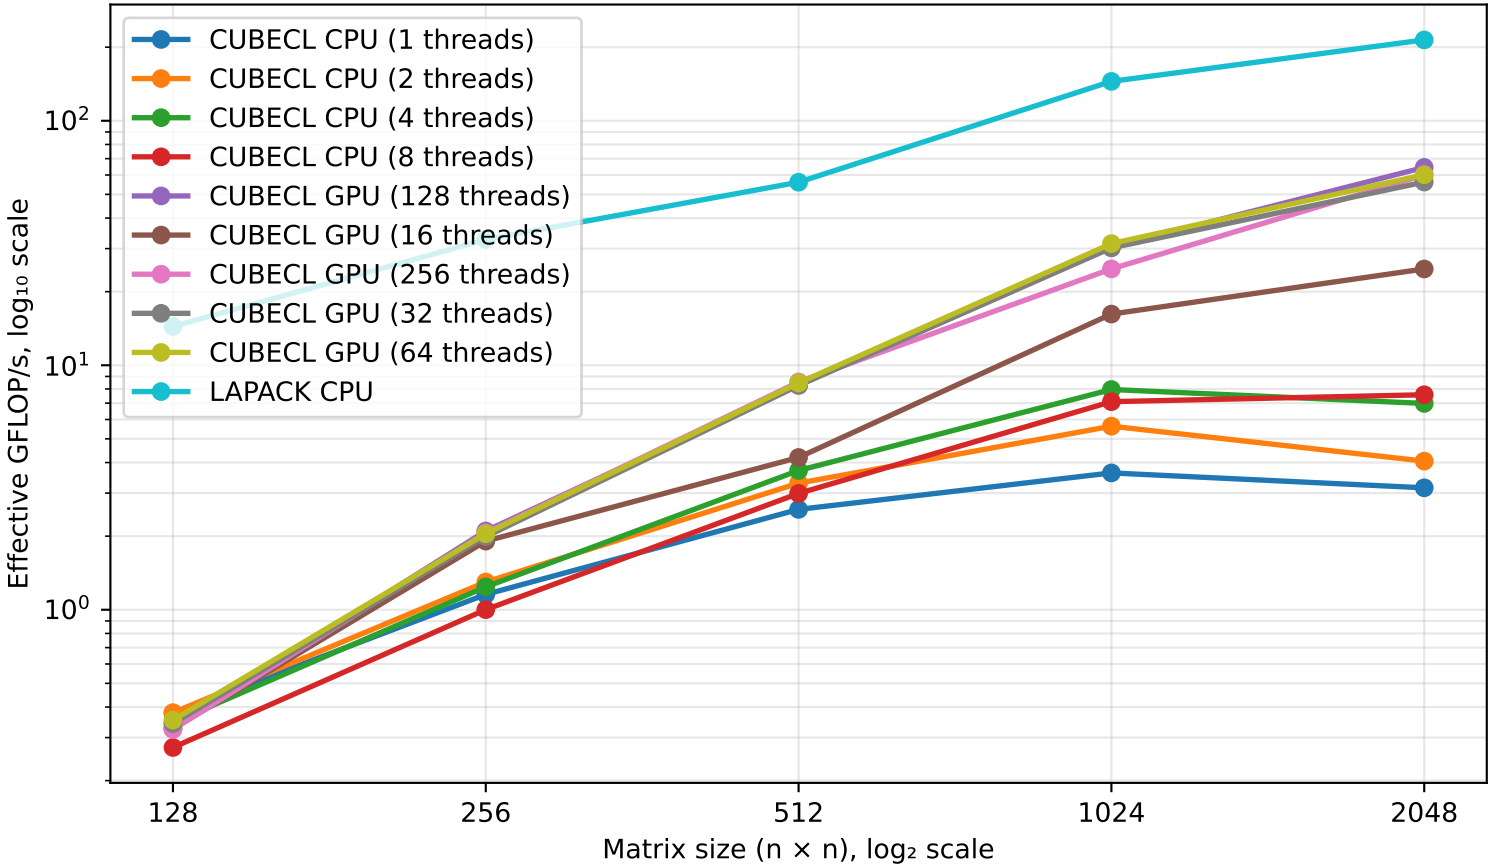
CPU vs GPU LU Throughput (log–log scale)
CUBECL CPU (1 threads)
CUBECL CPU (2 threads)
CUBECL CPU (4 threads)
CUBECL CPU (8 threads)
CUBECL GPU (128 threads)
CUBECL GPU (16 threads)
CUBECL GPU (256 threads)
CUBECL GPU (32 threads)
CUBECL GPU (64 threads)
LAPACK CPU
Effective GFLOP/s, log₁₀ scale
10^2
10^1
10^0
128
256
512
1024
2048
Matrix size (n × n), log₂ scale